



Министерство науки и высшего образования  
Российской Федерации  
Федеральное государственное бюджетное образовательное  
учреждение высшего образования  
"Московский государственный технический университет  
имени Н.Э. Баумана  
(национальный исследовательский университет)"  
(МГТУ им. Н.Э. Баумана)

---

ФАКУЛЬТЕТ ИНФОРМАТИКА, ИСКУССТВЕННЫЙ ИНТЕЛЛЕКТ И СИСТЕМЫ УПРАВЛЕНИЯ

КАФЕДРА \_\_\_\_\_СИСТЕМЫ ОБРАБОТКИ ИНФОРМАЦИИ И УПРАВЛЕНИЯ (ИУ5)\_\_\_\_\_

## ОТЧЕТ

### Лабораторная работа №7 «Алгоритмы Actor-Critic»

по курсу «Методы машинного обучения»

ИСПОЛНИТЕЛЬ:

группа ИУ5-21М

\_\_\_\_\_  
Савченко Г. А.

ФИО

\_\_\_\_\_  
подпись

"\_\_" \_\_\_\_\_ 2023 г.

ПРЕПОДАВАТЕЛЬ:

\_\_\_\_\_  
Гапанюк Ю.Е.

ФИО

\_\_\_\_\_  
подпись

"\_\_" \_\_\_\_\_ 2023 г.

---

Москва - 2023

## 1. Задание

- Реализуйте любой алгоритм семейства Actor-Critic для произвольной среды.

## 2. Текст программы

```
1 #!/usr/bin/env python
2
3 import gymnasium as gym
4 import numpy as np
5 from itertools import count
6 from collections import namedtuple
7
8 import torch
9 import torch.nn as nn
10 import torch.nn.functional as F
11 import torch.optim as optim
12 from torch.distributions import Categorical
13
14 # Cart Pole
15 CONST_ENV_NAME = 'Acrobot-v1'
16 env = gym.make(CONST_ENV_NAME)
17 GAMMA = 0.99
18 SavedAction = namedtuple('SavedAction', ['log_prob', 'value'])
19
20 class Policy(nn.Module):
21     def __init__(self):
22         super(Policy, self).__init__()
23         self.affine1 = nn.Linear(6, 128)
24
25         # actor's layer
26         self.action_head = nn.Linear(128, 3)
27
28         # critic's layer
29         self.value_head = nn.Linear(128, 1)
30
31         # action & reward buffer
32         self.saved_actions = []
33         self.rewards = []
34
35     def forward(self, x):
36         x = F.relu(self.affine1(x))
37
38         # actor: choses action to take from state s_t
39         # by returning probability of each action
40         action_prob = F.softmax(self.action_head(x), dim=-1)
41
42         # critic: evaluates being in the state s_t
43         state_values = self.value_head(x)
44
45         # return values for both actor and critic as a tuple of 2 values:
46         # 1. a list with the probability of each action over the action space
47         # 2. the value from state s_t
48         return action_prob, state_values
49
50 model = Policy()
51 optimizer = optim.AdamW(model.parameters(), lr=1e-3)
52 eps = np.finfo(np.float32).eps.item()
53
```

```

54 def select_action(state):
55     state = torch.from_numpy(state).float()
56     probs, state_value = model(state)
57
58     # create a categorical distribution over the list of probabilities of actions
59     m = Categorical(probs)
60
61     # and sample an action using the distribution
62     action = m.sample()
63
64     # save to action buffer
65     model.saved_actions.append(SavedAction(m.log_prob(action), state_value))
66
67     # the action to take (left or right)
68     return action.item()
69
70 def finish_episode():
71     """
72     Training code. Calculates actor and critic loss and performs backprop.
73     """
74     R = 0
75     saved_actions = model.saved_actions
76     policy_losses = [] # list to save actor (policy) loss
77     value_losses = [] # list to save critic (value) loss
78     returns = [] # list to save the true values
79
80     # calculate the true value using rewards returned from the environment
81     for r in model.rewards[::-1]:
82         # calculate the discounted value
83         R = r + GAMMA * R
84         returns.insert(0, R)
85
86     returns = torch.tensor(returns)
87     returns = (returns - returns.mean()) / (returns.std() + eps)
88
89     for (log_prob, value), R in zip(saved_actions, returns):
90         advantage = R - value.item()
91
92         # calculate actor (policy) loss
93         policy_losses.append(-log_prob * advantage)
94
95         # calculate critic (value) loss using L1 smooth loss
96         value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))
97
98     # reset gradients
99     optimizer.zero_grad()
100
101     # sum up all the values of policy_losses and value_losses
102     loss = torch.stack(policy_losses).sum() + torch.stack(value_losses).sum()
103
104     # perform backprop
105     loss.backward()
106     optimizer.step()
107
108     # reset rewards and action buffer
109     del model.rewards[:]
110     del model.saved_actions[:]
111
112 def main():
113     running_reward = -500
114

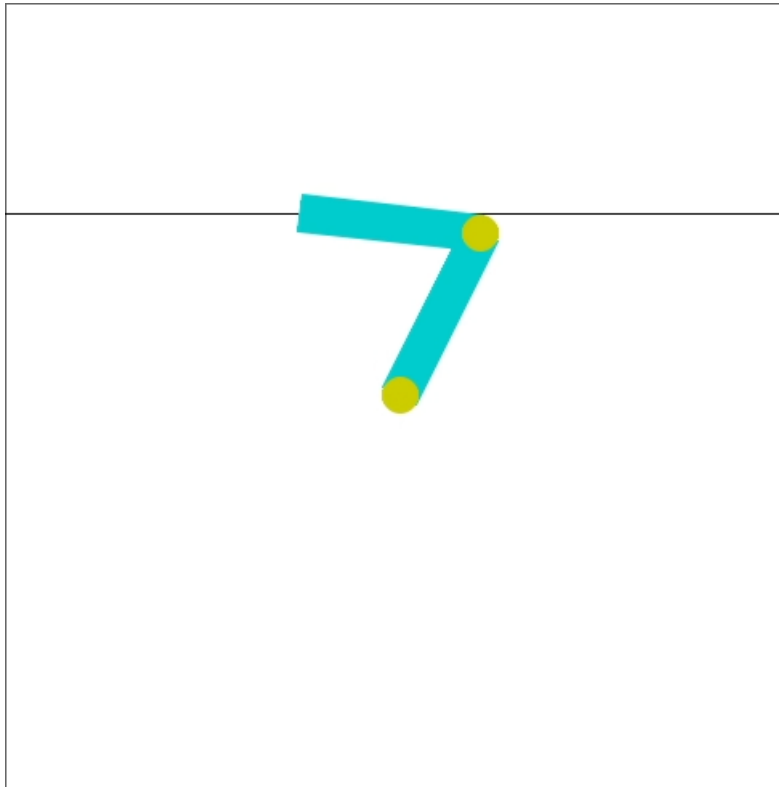
```

```

115 # run infinitely many episodes
116 for i_episode in count(1):
117     #print(running_reward)
118     # reset environment and episode reward
119     state, _ = env.reset()
120     ep_reward = 0
121
122     # for each episode, only run 9999 steps so that we don't
123     # infinite loop while learning
124     for t in range(1, 9999):
125         # select action from policy
126         action = select_action(state)
127
128         # take the action
129         state, reward, done, truncated, _ = env.step(action)
130
131         model.rewards.append(reward)
132         ep_reward += reward
133         if done or truncated:
134             break
135
136     print(ep_reward)
137     # update cumulative reward
138     running_reward = 0.05 * ep_reward + (1 - 0.05) * running_reward
139
140     # perform backprop
141     finish_episode()
142
143     # log results
144     if i_episode % 10 == 0:
145         print(f"Episode {i_episode}\tLast reward: {ep_reward:.2f}\tAverage reward:
146             {running_reward:.2f}")
147
148     # check if we have "solved" the cart pole problem
149     if running_reward > env.spec.reward_threshold*2:
150         print(f"Solved! Running reward is now {running_reward} and the last episode runs to {t}
151             time steps!")
152         break
153
154 env2 = gym.make(CONST_ENV_NAME, render_mode='human')
155
156 # reset environment and episode reward
157 state, _ = env2.reset()
158 ep_reward = 0
159
160 # for each episode, only run 9999 steps so that we don't
161 # infinite loop while learning
162 for t in range(1, 10000):
163     # select action from policy
164     action = select_action(state)
165     # take the action
166     state, reward, done, _, _ = env2.step(action)
167     model.rewards.append(reward)
168     ep_reward += reward
169     if done:
170         break
171
172 if __name__ == '__main__':
173     main()

```

### 3. Экранные формы с примерами выполнения программы



```
1 -500.0
2 -500.0
3 -500.0
4 -500.0
5 -500.0
6 -500.0
7 -500.0
8 -500.0
9 -500.0
10 -500.0
11 Episode 10      Last reward: -500.00    Average reward: -500.00
12 -500.0
13 -500.0
14 -500.0
15 -500.0
16 -500.0
17 -500.0
18 -500.0
19 -500.0
20 -500.0
21 -500.0
22 Episode 20      Last reward: -500.00    Average reward: -500.00
23 -500.0
24 -500.0
25 -500.0
26 -500.0
27 -500.0
28 -500.0
29 -500.0
30 -410.0
31 -500.0
32 -292.0
```

33 Episode 30	Last reward: -292.00	Average reward: -485.54
34 -232.0		
35 -335.0		
36 -500.0		
37 -500.0		
38 -500.0		
39 -500.0		
40 -500.0		
41 -500.0		
42 -461.0		
43 -500.0		
44 Episode 40	Last reward: -500.00	Average reward: -475.57
45 -500.0		
46 -500.0		
47 -500.0		
48 -424.0		
49 -362.0		
50 -391.0		
51 -350.0		
52 -500.0		
53 -257.0		
54 -500.0		
55 Episode 50	Last reward: -500.00	Average reward: -454.83
56 -500.0		
57 -384.0		
58 -500.0		
59 -481.0		
60 -500.0		
61 -292.0		
62 -500.0		
63 -342.0		
64 -428.0		
65 -468.0		
66 Episode 60	Last reward: -468.00	Average reward: -447.79
67 -500.0		
68 -329.0		
69 -472.0		
70 -500.0		
71 -412.0		
72 -299.0		
73 -500.0		
74 -500.0		
75 -284.0		
76 -455.0		
77 Episode 70	Last reward: -455.00	Average reward: -437.99
78 -241.0		
79 -333.0		
80 -500.0		
81 -315.0		
82 -339.0		
83 -292.0		
84 -258.0		
85 -408.0		
86 -273.0		
87 -297.0		
88 Episode 80	Last reward: -297.00	Average reward: -392.21
89 -451.0		
90 -387.0		
91 -268.0		
92 -317.0		
93 -255.0		

94	-223.0		
95	-352.0		
96	-267.0		
97	-258.0		
98	-242.0		
99	Episode 90	Last reward: -242.00	Average reward: -353.33
100	-208.0		
101	-273.0		
102	-283.0		
103	-165.0		
104	-231.0		
105	-195.0		
106	-237.0		
107	-306.0		
108	-213.0		
109	-267.0		
110	Episode 100	Last reward: -267.00	Average reward: -307.42
111	-242.0		
112	-149.0		
113	-236.0		
114	-258.0		
115	-196.0		
116	-204.0		
117	-152.0		
118	-366.0		
119	-251.0		
120	-285.0		
121	Episode 110	Last reward: -285.00	Average reward: -279.45
122	-210.0		
123	-143.0		
124	-185.0		
125	-231.0		
126	-142.0		
127	-253.0		
128	-251.0		
129	-322.0		
130	-160.0		
131	-162.0		
132	Episode 120	Last reward: -162.00	Average reward: -250.42
133	-313.0		
134	-205.0		
135	-186.0		
136	-162.0		
137	-186.0		
138	-186.0		
139	-217.0		
140	-124.0		
141	-171.0		
142	-208.0		
143	Episode 130	Last reward: -208.00	Average reward: -227.24
144	-214.0		
145	-207.0		
146	-149.0		
147	-125.0		
148	-245.0		
149	-204.0		
150	-175.0		
151	-243.0		
152	-135.0		
153	-459.0		
154	Episode 140	Last reward: -459.00	Average reward: -225.08

155 -175.0  
156 -163.0  
157 -212.0  
158 -177.0  
159 -166.0  
160 -157.0  
161 -175.0  
162 -140.0  
163 -156.0  
164 -197.0  
165 Episode 150      Last reward: -197.00      Average reward: -203.49  
166 -241.0  
167 -193.0  
168 -154.0  
169 -156.0  
170 Solved! Running reward is now -199.8982282342606 and the last episode runs to 157 time steps!